

Anjan Goswami, Ph.D.

San Francisco Bay Area, CA

AI Systems & Data Platform Executive

Full 4-page resume: smartinfer.com/resume/anjangoswami2026.pdf

goswami.anjan@gmail.com

linkedin.com/in/goswamianjan

smartinfer.com

Executive Summary

Senior Applied AI & ML Leader with 20+ years building and scaling AI systems across Microsoft, Amazon, Adobe, Salesforce, and others, from search/ranking and recommender systems to multimodal LLM products.

Specialized in multimodal LLM systems, evaluation, and enterprise-grade content generation.

Known for turning ambiguous research problems into shipped, high-quality AI products at scale, under latency, cost, safety, and compliance constraints.

Career Highlights

- **Microsoft PowerPoint Copilot:** Built and scaled the AI & Data org (1→16); architected agentic LLM systems for slide generation/editing and end-to-end evaluation and grounding; delivered **20%+ engagement** lift.
- **Adobe Document AI:** Designed a new high-performance PDF semantic parsing system with a privacy-preserving active learning loop; achieved **2× latency improvement** and improved accuracy from **45%→65%** while meeting strict throughput and compliance constraints.
- **Salesforce Einstein (Community & Service Cloud):** Built ML org from scratch; delivered feed ranking, recommender systems, trust/safety, and conversational AI driving **20% engagement lift** and improved enterprise adoption via vertical packages and instrumentation.
- **Walmart Commerce Search:** Led 80+ engineers/scientists in an AI re-architecture; drove **+23% revenue** and **+17% conversion**.
- **eBay & Amazon A9:** ML ranking and evaluation infrastructure for commerce discovery (eBay Cassini **+6% revenue**); global ranking infrastructure at Amazon A9.

Core Areas

Search & Personalization: ranking, retrieval, recommendation, marketplace discovery, experimentation and evaluation

Generative AI Systems: LLM applications, RAG, grounding, multimodal generation, agentic workflows, quality/eval frameworks

Data-Centric AI: active learning, privacy-preserving pipelines, synthetic data, long-tail mining, regression detection

Inference & Reliability: low-latency serving (cloud + edge), profiling, caching, distillation/compact models, production readiness

Professional Experience

Microsoft — Head of AI and Data → Senior Principal Scientist, PowerPoint (Office 365) Mar 2024 – Present

- Built and shipped agentic Copilot systems for slide generation/editing and on-brand grounded output; established evaluation, grounding, and regression frameworks.
- Built and scaled AI & Data org (1→16); partnered across product/engineering/research to drive multi-year AI strategy and model selection decisions.
- **Senior Principal scope (Dec 2025–Present):** Led controlled frontier model evaluations (diffusion LMs vs AR LLMs; infographic image models) to guide adoption decisions and future model/control strategy.

Adobe — Director of Machine Learning & Engineering

Jul 2019 – May 2023

- Led a global 40–60 person org across Acrobat, Sign, Stock, Creative Cloud, and Experience Cloud; restructured Document AI into product-aligned teams plus horizontal Data/Edge teams.

- Shipped data-centric document AI systems (semantic parsing, privacy-preserving active learning) and **Acrobat Assist** (RAG summarization/Q&A); delivered major latency/throughput and quality gains.
- Built CLIP-based semantic search for Adobe Stock and an audience discovery platform with **4×** cost reduction.

Salesforce — Director of Machine Learning & Engineering (Einstein)

Jul 2016 – Jul 2019

- Built ML org from scratch across Community & Service Cloud; delivered ranking, recommender systems, trust/safety, and conversational AI with **20%** engagement lift.
- Co-defined Service Cloud AI roadmap; introduced vertical packages and adoption instrumentation; established GDPR/HIPAA-aligned governance for customer-facing AI.

Walmart — Director of Search Science

Jun 2014 – Oct 2015

- Led 80+ engineers/scientists rebuilding commerce search (ranking, retrieval, type-ahead), delivering **+23% revenue** and **+17% conversion**.

Earlier: UpWork, eBay, Amazon A9, Microsoft

2005 – 2014

- UpWork: contractor ranking/recsys (+4% hiring) + experimentation. eBay: Cassini ranking (+6% revenue) + image search. Amazon A9: global ranking infra & evaluation.

Open Source & Selected Research

Thinker (Open Source): Unified LLM inference router (ThinkerQL, registry-based routing, observability). github.com/smartinfer/thinker

NeurIPS 2025 Workshop: VLM-SlideEval (multimodal evaluation). arXiv:2510.22045

Patents: 10+ in search relevance, ranking, and recommendations. smartinfer.com/patents

Education

Ph.D., Computer Science — University of California, Davis

M.S., Computer Science — The Ohio State University

M.Tech, Mechanical Engineering — IIT Kanpur